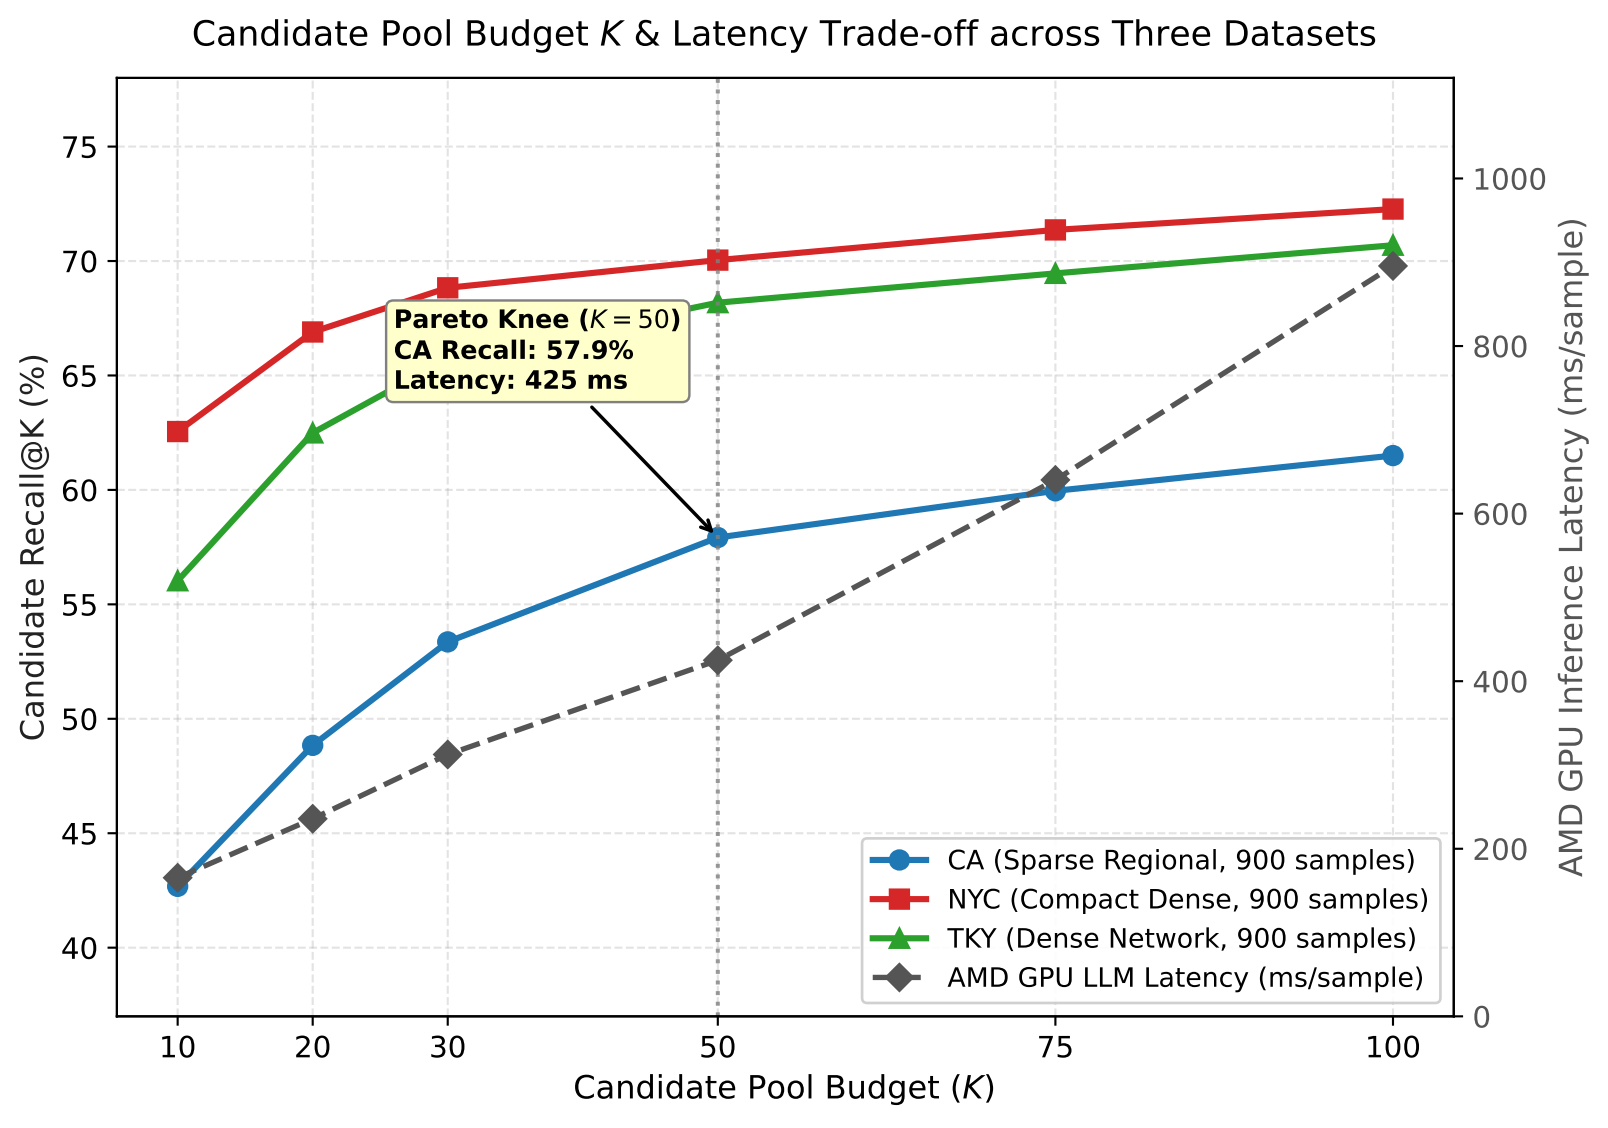
Candidate Pool Budget K & Latency Trade-off across Three Datasets
Pareto Knee ($K = 50$)
CA Recall: 57.9%
Latency: 425 ms
CA (Sparse Regional, 900 samples)
NYC (Compact Dense, 900 samples)
TKY (Dense Network, 900 samples)
AMD GPU LLM Latency (ms/sample)
Candidate Recall@K (%)
AMD GPU Inference Latency (ms/sample)
Candidate Pool Budget (K)
75
70
65
60
55
50
45
40
1000
800
600
400
200
0
10
20
30
50
75
100